

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

Load the Auto dataset in a DataFrame using Pandas

```
In [ ]: data = pd.read_csv('data/auto-dataset.csv')
print(data)
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130	3504	12.0	70	
1	15.0	8	350.0	165	3693	11.5	70	
2	18.0	8	318.0	150	3436	11.0	70	
3	16.0	8	304.0	150	3433	12.0	70	
4	17.0	8	302.0	140	3449	10.5	70	
..	...	...	...	...	...	...	...	
387	27.0	4	140.0	86	2790	15.6	82	
388	44.0	4	97.0	52	2130	24.6	82	
389	32.0	4	135.0	84	2295	11.6	82	
390	28.0	4	120.0	79	2625	18.6	82	
391	31.0	4	119.0	82	2720	19.4	82	

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino
..	...	...
387	1	ford mustang gl
388	2	vw pickup
389	1	dodge rampage
390	1	ford ranger
391	1	chevy s-10

[392 rows x 9 columns]

## 1. Scatterplots between features


Using the plot() function in matplotlib and create scatterplots between all the variables.

Is the relationship between those variables linear? Describe exactly four different connections between the variables.

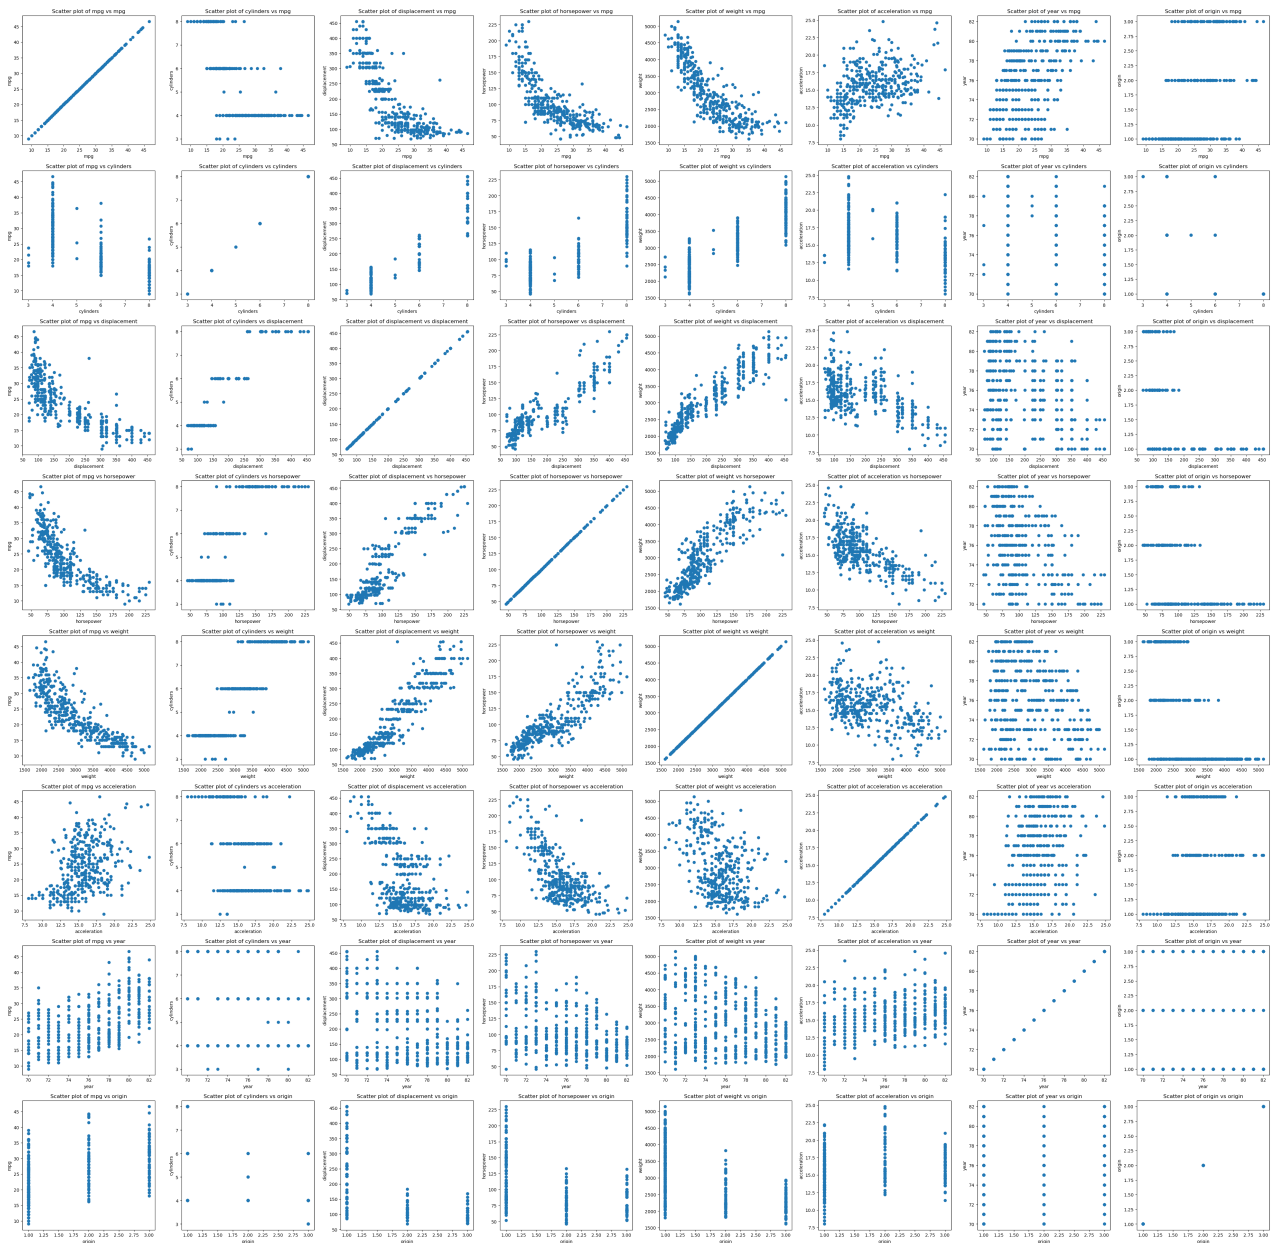
(Exclude the name variable, which is qualitative.)

```
In [ ]: feat1 = data['mpg']
feat2 = data['weight']
plt.figure()
plt.scatter(feat1, feat2, marker='x')
plt.xlabel('mpg')
plt.ylabel('weight')
plt.title('Scatter Plot of mpg vs weight')
#Similarly add scatter plots for every pair of features
```

Out[ ]: Text(0.5, 1.0, 'Scatter Plot of mpg vs weight')



```
In [ ]: # Plot the feature plots to observe the relationships between all the predictors
feats = ["mpg", "cylinders", "displacement", "horsepower", "weight", "acceleration", "y
figure, axis = plt.subplots(8, 8 , figsize=(50, 50), dpi=100)
for i in range(8):
    for j in range(8):
        index = (i ,j)
        plt.subplot2grid((8,8), index, rowspan = 1, colspan = 1)
        plt.scatter(data[feats[i]], data[feats[j]], label = '.')
        plt.title(f"Scatter plot of {feats[j]} vs {feats[i]}")
        plt.xlabel(feats[i])
        plt.ylabel(feats[j])
plt.show()
```



Describe exactly four different connections between the variables:

2/2

- 1) horsepower vs displacement appears to have a positive, roughly linear relationship.
- 2) horsepower vs acceleration appears to have a negative, roughly linear relationship (though starting to look like a banana).
- 3) mpg vs displacement, mpg vs horsepower, and mpg vs weight all show a negative, banana-shaped (quadratic, perhaps?) relationship.
- 4) year vs origin doesn't appear to have any sort of relationship.

## 2. Correlation

Detect the two variable pairs in the scatterplots that appear to be the most highly correlated and anti-correlated, respectively.

Justify your choice using the `np.corrcoef()` function.

Do the results from `np.corrcoef()` differ from what you see in the plot?

```
In [ ]: # Use np corrcoef to observe the correlation between predictors
corr_disp_wt = np.corrcoef(data["displacement"], data["weight"])
print(f"displacement vs weight: {corr_disp_wt[0][1]}")
# displacement vs weight looks highly correlated, and indeed has a correlation coefficient
corr_disp_hp = np.corrcoef(data["displacement"], data["horsepower"])
print(f"displacement vs horsepower: {corr_disp_hp[0][1]}")
corr_accel_year = np.corrcoef(data["acceleration"], data["year"])
print(f"acceleration vs year: {corr_accel_year[0][1]}")
corr_orig_year = np.corrcoef(data["origin"], data["year"])
print(f"origin vs year: {corr_orig_year[0][1]}")
# origin vs year looks pretty uncorrelated, and has a correlation of only 0.18

# But let's test it, just to be sure.
feats = ["mpg", "cylinders", "displacement", "horsepower", "weight", "acceleration", "year"]
min = 1
minners = ""
max = 0
maxers = ""
for i in range(8):
    for j in range(8):
        if i == j:
            continue
        score = abs(np.corrcoef(data[feats[i]], data[feats[j]]))
        if score[0][1] < min:
            min = score[0][1]
            minners = f"{feats[i]} vs {feats[j]}"
        if score[0][1] > max:
            max = score[0][1]
            maxers = f"{feats[i]} vs {feats[j]}"
print(f"Overall least correlated: {minners} @ {round(min,2)},\n Overall most correlated
```

```
displacement vs weight: 0.9329944040890104
displacement vs horsepower: 0.8972570018434686
acceleration vs year: 0.29031611333652
origin vs year: 0.18152771836633097
Overall least correlated: origin vs year @ 0.18,
Overall most correlated: cylinders vs displacement @ 0.95
```

Results:

displacement vs weight: 0.9329944040890104

displacement vs horsepower: 0.8972570018434686

acceleration vs year: 0.29031611333652

origin vs year: 0.18152771836633097

Overall least correlated: origin vs year @ 0.18,

Overall most correlated: cylinders vs displacement @ 0.95

### 3. Linear Regression

Perform simple linear regression with mpg as the response using the variables:

what about anti-cor?  
anti-correlation means negative correlation.

cylinders, displacement, horsepower, year

as features using the LinearRegression() function provided in the sklearn package.

Which predictors appear to have a statistically significant relationship to the outcome (use an appropriate measurement)?

How good are the resulting models (provide all four R2 values)?

```
In [ ]: from sklearn.linear_model import LinearRegression

linear_model = LinearRegression()

y = np.array(data['mpg'])
x = np.array(data['cylinders'])
print(x.shape, y.shape)
# Use the fit function and the score function in linear regression module of sklearn to
linear_model.fit(x.reshape(-1, 1), y)
linear_model.score(x.reshape(-1, 1), y)
```

(392,) (392,)

Out[ ]: 0.6046889889441246

```
In [ ]: # Fit the linear regression on mpg for the other features: cylinder, horsepower, year a
feats = ["cylinders", "displacement", "horsepower", "year"]
y = np.array(data['mpg'])
for feat in feats:
    linear_model = LinearRegression()
    x = np.array(data[feat])
    linear_model.fit(x.reshape(-1, 1), y)
    print(f"Score for mpg as a function of {feat}:", linear_model.score(x.reshape(-1, 1)
```

Score for mpg as a function of cylinders: 0.6046889889441246

Score for mpg as a function of displacement: 0.6482294003193044

Score for mpg as a function of horsepower: 0.6059482578894348

Score for mpg as a function of year: 0.33702781330962295

Results:

Score for mpg as a function of cylinders: 0.6046889889441246

Score for mpg as a function of displacement: 0.6482294003193044

Score for mpg as a function of horsepower: 0.6059482578894348

Score for mpg as a function of year: 0.33702781330962295

4/4

mpg as a function of year seems to be the least correlated, with an  $R^2$  of only 0.34

$R^2$  shows the goodness of fit

Out of a best possible  $R^2$  value of 1.0, the other variables being between 0.60 and 0.65 does not seem to be all that high of an impact.

## 4. Multiple Linear Regression

Use the `LinearRegression()` function to perform one multiple linear regression with mpg as the response and all other variables except name as the predictors.

2/4

Use the `score()` and `get_params()` functions to print the results.

Compare the full model to those generated in 5.3:

Comment on the goodness of fit?

(a) How is the model fit (using  $R^2$ )?

Comment on the sign of coeffs?

(b) What can you observe in the different models concerning the significance of the relationship between response and individual predictors?

What does the sign of the coefficient (i.e. of the estimate) tell you about the relationship between the predictor and the response?

Provide an example from your fitted model.

In [ ]:

```
#Multiple Linear Regression
#Create a 2D matrix with each column representing all features
#Look up np.concatenate or np.stack
feats = ["mpg", "cylinders", "displacement", "horsepower", "weight", "acceleration", "y"]
mpg_array = np.array(data["mpg"])
cylinders_array = np.array(data["cylinders"])
displacement_array = np.array(data["displacement"])
horsepower_array = np.array(data["horsepower"])
weight_array = np.array(data["weight"])
acceleration_array = np.array(data["acceleration"])
year_array = np.array(data["year"])
origin_array = np.array(data["origin"])

combined_arrays = np.stack((cylinders_array, displacement_array, horsepower_array, weight_array, acceleration_array, year_array, origin_array))

#Print the predictions of the your fitted model and the weights of Linear Regression (L)
multiple_linear_model = LinearRegression()
multiple_linear_model.fit(combined_arrays, mpg_array)
mlm_score = multiple_linear_model.score(combined_arrays, mpg_array)
print(mlm_score)
#params = multiple_linear_model.get_params()
#print(params)
coefs = multiple_linear_model.coef_
print(coefs)
predictions = multiple_linear_model.predict(combined_arrays)
print(predictions)
```

0.8214780764810599

[-0.49337632 0.01989564 -0.01695114 -0.00647404 0.08057584 0.75077268  
 1.4261405 ]

[15.00095865 13.99929917 15.24044696 15.06190592 14.96717762 10.69562338  
10.6553509 10.69318951 10.21140958 13.11319398 15.29186157 14.14690266  
14.64696189 18.88015369 24.13062472 19.03847473 19.40643664 20.88068005  
25.45738302 27.13648841 21.03963303 22.25408365 22.7074409 23.25050717  
20.3005362 7.58306997 8.41074023 8.30037285 6.44696729 26.20815569  
23.39053912 25.73308381 21.46783861 16.23364804 17.52793963 17.90615253  
17.43525382 11.44975337 10.58385513 12.14999889 11.87950268 6.98211816  
8.88321853 6.24768971 19.73241559 23.0454129 17.79193175 18.88084577  
23.26421626 25.13116053 25.49443983 25.22718096 28.75457004 29.66193221  
27.53555535 25.15803406 26.28101808 24.47954013 26.03152794 23.5313529  
24.17758427 11.77971323 11.88636515 12.41807558 13.02233473 14.97586698  
10.21781246 10.59400653 10.80705681 11.4328463 25.38717749 13.63215327]

12.81807391	11.44132342	12.83385802	20.40479902	24.02910196	20.91375517
25.73501085	23.04318302	26.06966453	24.82972756	24.08232262	27.31430135
13.56803385	15.72663966	14.80166104	13.75354471	15.40597996	9.11316466
12.54946701	12.17342553	12.58907167	10.40842679	9.13895663	15.41209672
19.87451504	19.5611008	21.19768398	21.30747992	20.8714236	28.68457937
9.25995356	9.44005829	10.26512597	10.8598469	22.12705892	27.10765986
24.63256448	26.37596114	27.59026112	24.56264013	22.44813226	25.47613389
14.19310096	12.19610795	28.47323412	26.78124528	23.60437753	21.65696035
17.7323046	22.91659991	22.85694927	16.14480938	20.3806236	22.23331457
19.85580312	29.98624525	24.32417673	30.76017502	24.21820401	16.97485381
18.00971582	17.44005836	13.77224574	10.9854969	11.83497896	10.71579784
13.01240031	26.72394641	28.19392676	26.1032573	31.92807295	29.82441039
25.72165832	27.33843339	26.84143838	26.48161957	26.9854206	27.99467507
20.57949571	19.6449369	20.78200302	22.43502716	12.35099671	13.45751462
12.4008917	11.98140149	16.6846717	17.01074549	18.20660808	17.44824662
21.82068634	20.151559	20.84109951	29.28703926	24.20538532	22.89992468
24.71528484	26.02804741	27.4930715	27.01110186	21.16309705	29.16015946
21.22722108	24.41349693	23.09011142	22.81674455	24.22188715	32.09569464
26.68698266	28.64327193	25.06814935	26.73078129	28.29854643	14.77382297
15.02480599	16.80340145	15.46949369	21.39786251	20.96109442	22.83057359
22.84206873	28.71669964	27.99652467	29.9270473	32.83840974	18.8770368
20.51175467	19.12289695	22.59151822	30.61330685	31.13619763	30.17362697
25.34774195	22.44828555	16.659432	22.20933012	24.71142766	17.65541354
13.85034576	16.53964233	17.27834413	17.9215907	31.94430733	28.1369806
31.74200818	27.20228117	32.16218709	17.56835655	16.62475777	16.24043964
15.26959078	20.69904022	21.06110075	19.76376439	21.07958092	16.55117173
16.10614612	15.79051915	15.63121352	30.68623095	25.10657282	30.35729138
24.75754377	29.01707404	28.4340407	32.16059618	29.03548977	26.11364564
26.21144222	26.59475552	32.07896737	31.132481	33.0293463	32.24874085
34.10835103	21.59297517	19.6576677	20.39227537	21.28505309	23.21736023
24.4235579	25.71962878	21.76905783	23.50212832	22.01457196	23.8163016
20.4903419	22.04686279	21.32122206	20.38330848	22.62637987	17.46441064
28.96950257	29.27303384	30.66423609	28.34340913	29.61271129	25.69268713
25.29260528	30.0005767	25.54720814	23.12528937	25.84869065	21.41895534
31.10502466	31.95932201	23.43759872	25.20586151	25.52290817	23.97740445
22.7551566	19.81680074	20.35742826	19.67460936	20.03931365	16.88153561
19.05450791	20.63469186	19.73125191	32.2040919	33.44524117	30.90142267
26.51946521	23.57987219	20.5695695	25.91921953	22.86655646	29.26841024
29.72909134	33.45631692	31.00700967	27.02739892	26.16141866	25.61056544
27.53982752	31.68775723	34.67584782	30.58316132	34.00784694	27.76666717
26.46762239	25.95267609	23.87820651	31.43773481	29.98650295	31.21749611
31.40718811	32.58673398	33.53957275	26.65283567	33.6244257	32.94922355
31.49186438	27.0054564	26.20710489	34.95803903	33.50598861	33.73009694
27.34555857	30.67926574	29.68930984	32.60087719	29.4534759	28.889473
28.78930104	27.01084268	29.90763811	36.4874435	32.85555202	36.39650161
34.75856959	35.2653094	34.74098289	34.99513155	30.92848188	31.96064606
30.15443279	32.27360779	33.60922273	32.91613431	30.83443461	31.44759457
26.65470916	26.20406605	28.52898099	27.92594196	23.95554276	23.67201668
26.11473697	24.000706	29.24857498	28.94140762	30.47920275	29.21969186
30.00182981	29.02005581	27.79675446	34.472563	35.66420933	35.97193622
32.2458033	32.19737296	34.73003366	34.40042321	34.43949485	35.8117604
35.86816348	35.7142301	26.80267915	28.43316736	29.77472282	28.35862054
31.75418253	30.76357903	27.5717334	28.31955391	34.46457181	31.13632611
29.35024372	28.72892119]				

## 5. Residuals


Use the `plot()` function to produce the residual versus fitted plot of the multiple linear regression fit.

Identify the residual plot.

Does the residual plot suggest any non-linearity in the data (provide an explanation)?

```
In [ ]: #Make a scatter plot of the residual vs the predictions of your linear regression model
#predictions = multiple_linear_model.predict(combined_arrays)
residuals = (data['mpg'] - predictions)
plt.figure()
plt.scatter(predictions, residuals, marker='.')
plt.xlabel('Fitted values')
plt.ylabel('Residuals')
plt.title("Scatter plot of residuals vs fitted values")
plt.axhline(y = np.nanmean(residuals), color = "red")
```

```
Out[ ]: <matplotlib.lines.Line2D at 0x19b8639d7f0>
```



2/2

The residual plot shows non-linearity (banana in space), indicating a potential quadratic term in the model.

The residuals also display evidence of heteroscedasticity, increasing variance as the fitted value increased.